

GENESIS OF SDN

3

We have now seen a number of the reasons that have led to the emergence of SDN. This included the fact that there were too many disparate control plane protocols attempting to solve many different networking problems in a distributed fashion, which has resulted in relative confusion and inertia in network switching evolution as compared to compute and storage technologies. Overall, the system that has evolved is far too closed, and there is a strong desire to migrate to a model closer to that of open source development and open source initiatives (such as Linux with respect to computer operating systems). We have also noted that for a variety of reasons the *network equipment manufacturers* (NEMs) are bogged down in their current technologies and are unable to innovate at the rate required by the modern data center. This is partly attributable to the fact that the current business model for networking equipment is a very high margin business and it is not obviously in the incumbents' interest for the status quo to change drastically [1]. Even the standards bodies generally move too slowly [2]. Thus, the need for an SDN-like transformative technology has become apparent to many users of network equipment, especially in large data centers. In this chapter we will explore how the SDN movement began. In order to understand the base from which these SDN pioneers started, we begin with a brief review of the evolution of networking technology up to the advent of SDN.

3.1 THE EVOLUTION OF NETWORKING TECHNOLOGY

In the earliest days of computing there was no networking—each individual computer occupied an entire room or even an entire floor. As mainframes emerged and began to find their way into our technological consciousness, these computing powerhouses still operated as islands, and any sharing of data took place using physical media such as magnetic tapes.

3.1.1 MAINFRAME NETWORKING: REMOTE TERMINALS

Even in the age of mainframes, remote connectivity to the mainframe was needed. This was provided in the form of remote terminal controllers and card readers, which operated as subservient devices, known as peripherals, with control residing entirely in the central mainframe. Network connections in this case were simple point-to-point or point-to-multipoint links. Communication was solely between a few connected entities, in this case the mainframe and a small set of remote terminal controllers or card readers. This communication was controlled by the central mainframe.

3.1.2 PEER-TO-PEER POINT-TO-POINT CONNECTIONS

As computer technology began to move from solely mainframes to the addition of mini-computers, these machines had a greater need to share information in a quick and efficient manner. Computer manufacturers began to create protocols for sharing data between two peer machines. The network in this case was also point-to-point, although the nature of the connection was peer-to-peer, in that the two machines (e.g., mini-computers) would communicate with each other and share data as relative equals, at least compared to the earlier mainframe-to-terminal-controller type of connections.

Of course, in these point-to-point connections, the network was trivial, with only the two parties communicating with each other. Control for this communication resided not in any networking device, but in the individual computers participating in this one-to-one communication.

3.1.3 LOCAL AREA NETWORKS

Eventually, with further evolution of computing toward smaller, independent systems, the need arose for a way to connect these devices in order to allow them to share information and collaborate in a manner that wasn't required when everything ran on one large mainframe or even on a powerful minicomputer. Hence, *Local Area Networking* (LAN) technology arrived, with various battles being fought between technologies (e.g., Ethernet/IEEE 802.3 versus Token Ring/IEEE 802.5). Notably, these early LANs were running on shared media, so all traffic was seen by every device attached to the network.

IEEE 802.3 emerged as the more popular of these technologies. It uses *Carrier-Sense Multiple-Access/Collision Detect* (CSMA/CD) technology, which exhibits poor aggregate throughput when the number of active devices reached a certain level. The exact number of devices where this would occur was dependent on the amount of data attempted to be transmitted by each. This decrease in performance resulted from the backing off and waiting to transmit done by each device when a collision occurs, as stipulated by CSMA/CD. The number of collisions reaches a critical point as a function of the number of nodes and their respective transmission activity. Once this critical point is reached the network becomes very inefficient, with too much time spent either in the midst of or recovering from collisions.

These flat, shared-media networks were quite simple, and the *repeater* devices provided physical extensions to the shared media by merely forwarding all frames to the extended medium. The greatest impact of these early devices was a new topology created by wire concentration. This made layer two networks more deployable and reliable than the former method of snaking a coaxial cable from one node to the next. This early networking required minimum control plane intelligence, if any. Simple repeaters basically just did one thing: forward to everybody. More advanced, *managed* repeaters did have limited control planes where segmentation and error monitoring occurred. Managed repeaters did perform functions such as removing erroneous traffic from the network and also isolated ports that were causing problems. Segmentable repeaters were able to split themselves into multiple repeaters via configuration. Different groups of ports could be configured to reside in different collision domains. These separate collision domains would be connected by bridges. This feature provided more control over the size of the collision domains to optimize performance.

DISCUSSION QUESTION:

Explain why shared medium networks often cease to work efficiently above a certain level of utilization.

3.1.4 BRIDGED NETWORKS

Eventually these shared-media networks needed to scale in physical extent as well as number of nodes. As explained in the previous section, shared-media networks do not scale well as the number of hosts grows. It became desirable to split the shared-media network into separate segments. In so doing, and since not all the nodes are transmitting all the time, *spatial reuse* occurs and the aggregate bandwidth available actually increases due to the locality of transmissions in each segment. The first devices to perform this functionality were called *bridges*—forerunners of today's switches, but much simpler. They typically had only two ports, connecting two shared domains. These bridges possessed an incipient control plane in that they were able to actually learn the location and MAC address of all devices, and created forwarding tables which allowed them to make decisions about what to do with incoming packets. As we mentioned in the previous chapter, these forwarding tables were implemented entirely in software.

Furthermore, these bridges were implemented in such a way that each device was able to operate independently and autonomously without requiring any centralized intelligence. The goal was to facilitate expansion of the network without a lot of coordination or interruption across all the devices in the network. One of the first manifestations of this distributed intelligence paradigm was the *Spanning Tree Protocol*, (STP) which allowed a group of bridges to interact with each other and converge on a topology decision (the spanning tree), which eliminated loops and superimposed a hierarchical loop-free structure on the network.

The important point here is that this greater scale in terms of the number of nodes as well as physical extent drove the need for this new model of individually autonomous devices, with distributed protocols implementing the required control functionality. If we translate this period's realities into today's terminology, there was no centralized controller, merely a collector of statistics. *Policies*, if indeed one can use that term, were administered by setting specific parameters on each device in the network. We need to keep in mind that at the time networks in question were small, and this solution was entirely acceptable.

3.1.5 ROUTED NETWORKS

In the same manner that bridged and switched networks dealt with layer two domains with distributed protocols and intelligence, similar strategies were employed for layer three routing. Routers were directly connected locally to layer two domains and interconnected over large distances with point-to-point WAN links. Distributed routing protocols were developed to allow groups of interconnected routers to share information about those networks to which they were directly connected. By sharing this information amongst all the routers, each was able to construct a routing table allowing it to route packets to remote layer three IP subnets using the optimal forwarding ports. This was another application of autonomous devices utilizing distributed protocols in order to allow each to make appropriate forwarding decisions.

This has led to the current state of affairs, with networking intelligence distributed in the networking devices themselves. During this evolution, however, the growth of network size and complexity was unrelenting. The size of MAC forwarding tables grew, control plane protocols became more complex, and network overlays and tunneling technology became more prevalent. Making major changes to these implementations was a continual challenge. Since the devices were designed to operate independently,

centrally administered large-scale upgrades were challenging. In addition, the fact that the actual control plane implementations were from many different sources, and not perfectly matched, created a sort of *lowest common denominator* effect, where only those features that were perfectly aligned between the varied implementations could truly be relied upon. In short, existing solutions were not scaling well with this growth and complexity. This led network engineers and researchers to question whether this evolution was headed in the right direction. We describe some of the innovations and research that this led to in the following section.

3.2 FORERUNNERS OF SDN

Prior to OpenFlow, and certainly prior to the birth of the term Software Defined Networking, forward-thinking researchers and technologists were considering fundamental changes to today's world of autonomous, independent devices and distributed networking intelligence and control. This section considers some of those early explorations of SDN-like technology. There is a steady progression of ideas around advancing networking technology toward what we now know as SDN. [Table 3.1](#) shows this progression, which will be discussed in more detail in the following subsections. The timeframes shown in the table represent approximate time periods when these respective technologies were developed or applied in the manner described.

Table 3.1 Precursors of SDN	
Project	Description
Open Signaling	Separating the forwarding and control planes in ATM switching (1999, see Section 3.2.1)
Active Networking	Separate control and programmable switches (late 1990s, see Section 3.2.1)
DCAN	Separating the forwarding and control planes in ATM switching (1997, see Section 3.2.1)
IP Switching	Control layer two switches as a layer three routing fabric (late 1990s, see Section 3.2.1)
MPLS	Separate control software establishes semi-static forwarding paths for flows in traditional routers (late 1990s, see Section 3.2.1)
RADIUS, COPS	Use of admission control to dynamically provision policy (2010, see Section 3.2.2)
Orchestration	Use of SNMP and CLI to help automate configuration of networking equipment (2008, see Section 3.2.3)
Virtualization Manager	Use of plug-ins to perform network reconfiguration to support server virtualization (2011, see Section 3.2.4)
ForCES	Separating the forwarding and control planes (2003, see Section 3.2.5)
4D	Control plane intelligence located in a centralized system (2005, see Section 3.2.6)
Ethane	Complete enterprise and network access and control using separate forwarding and control planes, and utilizing a centralized controller (2007, see Section 3.2.7)

3.2.1 EARLY EFFORTS

Good surveys of early programmable networks that were stepping-stones on the path to SDN can be found in [3,4]. Some of the earliest work in programmable networks began not with Internet routers and switches, but with ATM switches. Two notable examples were *Devolved Control of ATM Networks* (DCAN) and *Open Signaling*.

As its name indicates, DCAN [5] prescribed the separation of the control and management of the ATM switches from the switches themselves. This control would be assumed by an external device that is similar to the role of the controller in SDN networks.

Open Signaling [6] proposed a set of open, programmable interfaces to the ATM switching hardware. The key concept was to separate the control software from the switching hardware. This work led to the IETF effort which resulted in the creation of the *General Switch Management Protocol* (GSMP) [7]. In GSMP, a centralized controller is able to establish and release connections on an ATM switch, as well as a multitude of other functions that may otherwise be achieved via distributed protocols on a traditional router. Tag switching was Cisco's version of label switching, and both of these terms refer to the technology that ultimately became known as *Multiprotocol Label Switching* (MPLS). Indeed, MPLS and related technologies are a deviation from the autonomous, distributed forwarding decisions characteristic of the traditional Internet router and in that sense were a small step toward a more SDN-like Internet switching paradigm. In the late 1990s, Ipsilon Networks utilized the GSMP protocol to set up and tear down ATM connections internally in their IP Switch product. The Ipsilon IP Switch [8] presented normal Internet router interfaces externally, but its internal switching fabric could utilize ATM switches for persistent flows. Flows were defined as a relatively persistent set of packets between the same two endpoints, where an endpoint was determined by IP address and TCP/UDP port number. Since there was some overhead in establishing the ATM connection that would carry that flow, this additional effort would only be expended if the IP Switch believed that a relatively large number of packets would be exchanged between the endpoints in a short period of time. This definition of flow is somewhat consistent with the notion of flow within SDN networks.

The *Active Networking* project [9,10] also included the concept of switches that could be programmed by out-of-band management protocols. In the case of Active Networking, the switching hardware was not ATM switches but Internet routers. Active Networking also included a very novel proposal for small downloadable programs called *capsules* that would travel in packets to the routers and could reprogram the router's behavior on the fly. These programs could be so fine-grained as to prescribe the forwarding rules for a single packet, even possibly the payload of the packet that contained the program. The concept of Active Networking was embraced by a number of projects [11–13].

3.2.2 NETWORK ACCESS CONTROL

Network Access Control (NAC) products control access to a network based on policies established by the network administration. The most basic control is to determine whether or not to admit the user onto the network. This is usually accomplished by some exchange of credentials between the user and the network. If the user is admitted, his rights of access will also be determined and restricted in accordance with those policies. There were early efforts where network policy beyond admission control was dynamically provisioned via NAC methods, such as *Remote Authentication Dial In User*

Service (RADIUS) [14] and *Common Open Policy Service* (COPS) [15]. As RADIUS was far more widely implemented, we discuss it in more detail here.

RADIUS has been used to provide the automatic reconfiguration of the network. RADIUS can be viewed as a simple, proactive, policy-based precursor to SDN. The idea is that via RADIUS, networking attributes would change based on the identity of the compute resource that had just appeared and required connectivity to the network. While RADIUS was originally designed for the *Authentication, Authorization and Accounting* (AAA) related to granting a user access to a network, that original paradigm maps well to our network reconfiguration problem. The identity of the resource connecting to the network serves to identify the resource to the RADIUS database, and the authorization attributes returned from that database could be used to change the networking attributes described above. Solutions of this nature achieved automatic reconfiguration of the network, but never gained the full trust of IT administrators and, thus, never became mainstream. While this RADIUS solution did an adequate job of automatically reconfiguring the edge of the network, the static, manually configured core of the network remained the same. This problem still awaited a solution.

Fig. 3.1 shows an overview of the layout and operation of a RADIUS-based solution for automating the process of creating network connectivity for a virtual server. In the figure, the process would be that a VM is moved from physical server A to physical server B, and as a result, the RADIUS server becomes aware of the presence of this VM at a new location. The RADIUS server is then able to automatically configure the network based on this information, using standard RADIUS mechanisms.

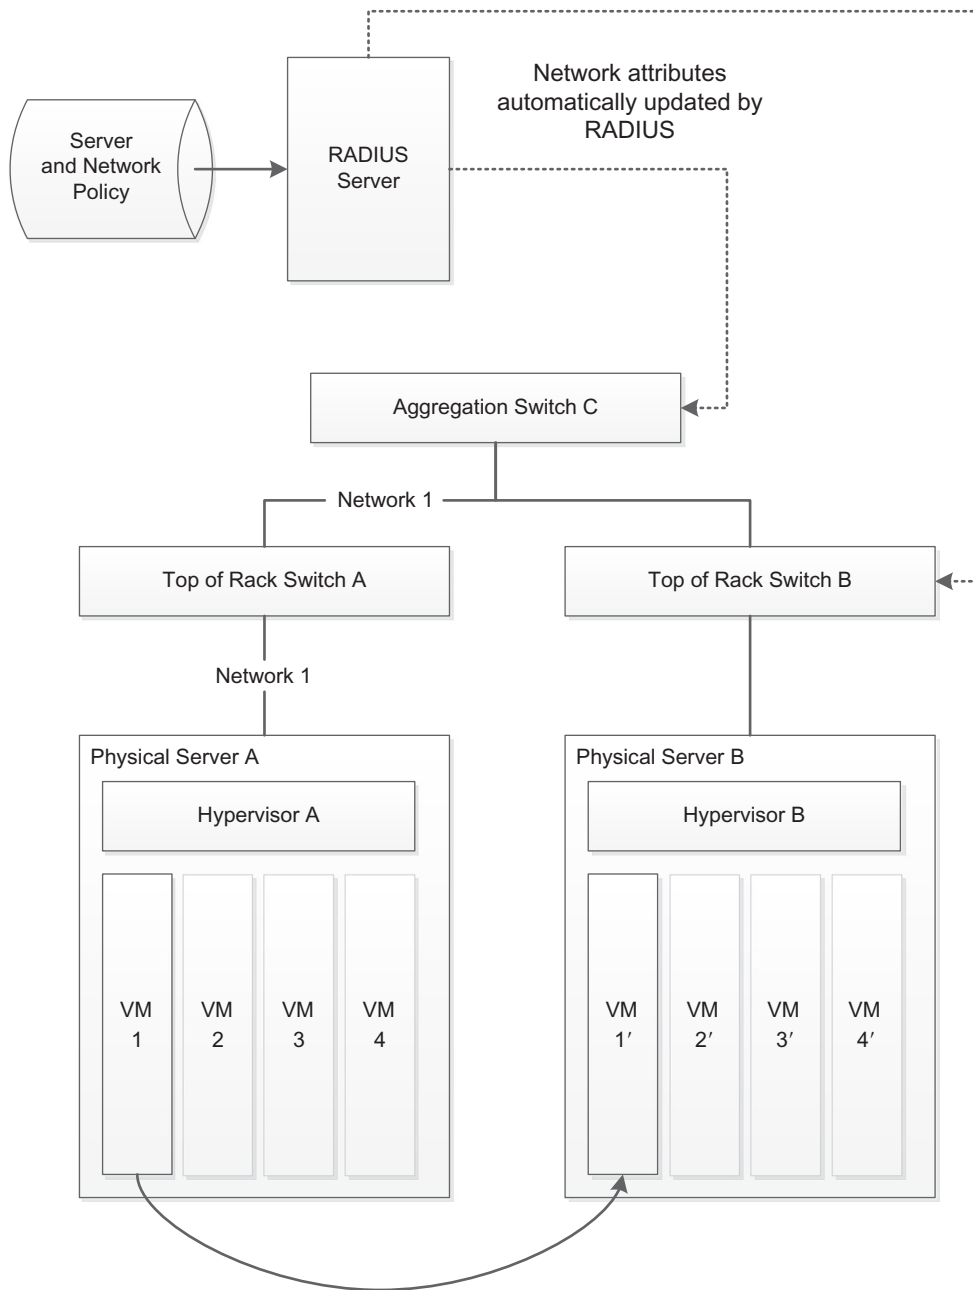
3.2.3 ORCHESTRATION

Early attempts at automation involved applications which were commonly labeled *Orchestrators* [16]. Just as a conductor can make a harmonious whole out of the divergent instruments of an orchestra, such applications could take a generalized command or goal, and apply it across a wide range of often heterogeneous devices. These orchestration applications would typically utilize common device *Application Programming Interfaces* (APIs), such as the *Command Line Interface* (CLI) or *Simple Network Management Protocol* (SNMP).¹

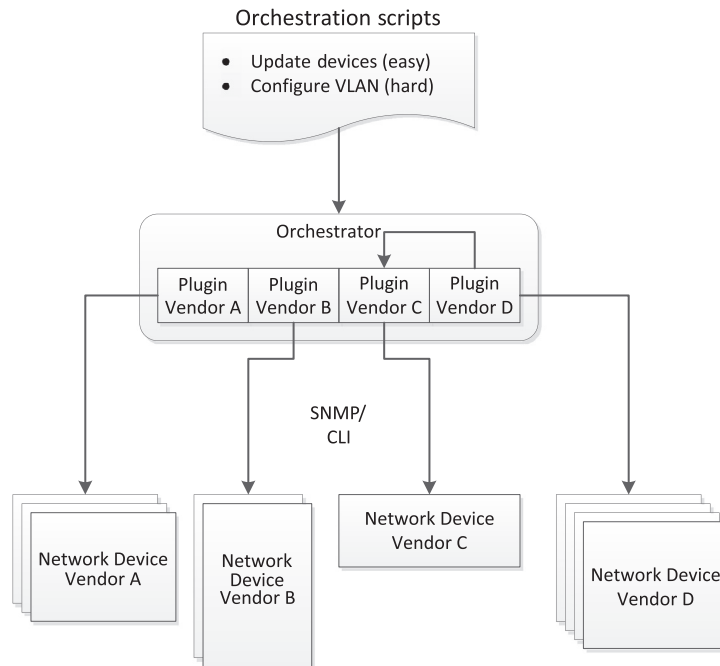
Fig. 3.2 shows an overview of the layout and operation of an orchestration management solution for automating the process of creating network connectivity for a virtual server. The figure shows SNMP/CLI plug-ins for each vendor's specific type of equipment; an orchestration solution can then have certain higher-level policies which are in turn executed at lower levels by the appropriate plug-ins. The vendor-specific plug-ins are used to convert the higher-level policy requests into the corresponding native SNMP or CLI request specific to each vendor.

Such orchestration solutions alleviated the task of updating device configurations. But since they were really limited to providing a more convenient interface to existing capabilities, and since no capability existed in the legacy equipment for network-wide coordination of more complex policies such as security and virtual network management, tasks such as configuring VLANs remained hard.

¹The reader should note that while in this work we classify CLI and SNMP as APIs, we acknowledge that they are much more primitive than most of the APIs we discuss later in the book and there is not universal consensus that they should even be considered APIs.

**FIG. 3.1**

Early attempts at SDN: RADIUS.

**FIG. 3.2**

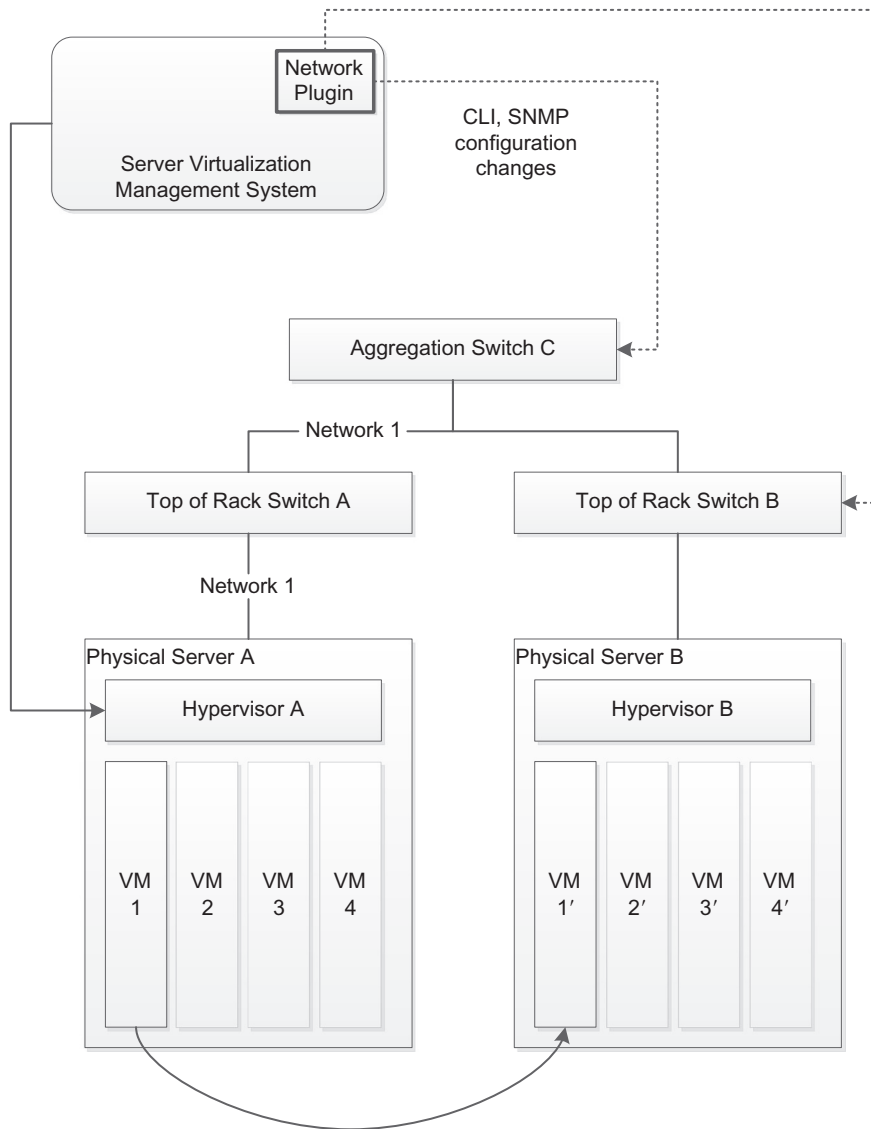
Early attempts at SDN: Orchestration.

Consequently, orchestration management applications continued to be useful primarily for tasks such as software and firmware updates, but not for tasks which would be necessary in order to automate today's data centers.

3.2.4 VIRTUALIZATION MANAGER NETWORK PLUG-INS

The concept of *Virtualization Manager Network Plug-ins* builds upon the notion of orchestration and attempts to automate the network updates that are required in a virtualization environment. Tools specifically targeted at the data center would often involve virtualization manager plug-ins (e.g., plug-ins for VMware's vCenter [17]) which would be configured to take action in the event of a server change, such as a vMotion [18]. The plug-in would then take the appropriate actions on the networking devices they controlled in order to make the network follow the server and storage changes with changes of its own. Generally the mechanism for making changes to the network devices would be SNMP or CLI commands. These plug-ins can be made to work, but, since they must use the static configuration capabilities of SNMP and CLI, they suffer from being difficult to manage and prone to error.

Fig. 3.3 shows an overview of the layout and operation of a virtualization manager plug-in solution for automating the process of creating network connectivity for a virtual server. This figure reflects a

**FIG. 3.3**

Early attempts at SDN: Plug-ins.

similar use of SNMP/CLI plug-ins to that which we saw in [Fig. 3.2](#). The most notable difference is that this application starts from a base that supports virtualization. The nature of the reconfiguration managed by this application is oriented to support the networking of virtual machines via VLANs and tunnels.

3.2.5 FORCES: SEPARATION OF FORWARDING AND CONTROL PLANES

The *Forwarding and Control Element Separation* (ForCES) [19] work produced in the IETF began around 2003. ForCES was one of the original proposals recommending the decoupling of forwarding and control planes as well as a standard interface for communication between them. The general idea of ForCES was to provide simple hardware-based forwarding entities at the foundation of a network device, and software-based control elements above. These simple hardware forwarders were constructed using cell switching or tag switching technology. The software-based control had responsibility for the broader tasks often involving coordination between multiple network devices (e.g., Border Gateway Protocol routing updates).

The functional components of ForCES are as follows:

- **Forwarding Element:** The *Forwarding Element* (FE) would be typically implemented in hardware and located in the network. The FE is responsible for enforcement of the forwarding and filtering rules that it receives from the *controller*.
- **Control Element:** The *Control Element* (CE) is concerned with the coordination between the individual devices in the network, and for communication of forwarding and routing information to the FEs below.
- **Network Element:** The *Network Element* (NE) is the actual network device which consists of one or more FEs and one or more CEs.
- **ForCES Protocol:** The ForCES protocol is used to communicate information back and forth between FEs and CEs.

ForCES proposes the separation of the forwarding plane from the control plane, and it suggests two different embodiments of this architecture. In one of these embodiments, both the forwarding and control elements are located within the networking device. The other embodiment speculates that it would be possible to actually move the control element(s) off of the device and to locate them on an entirely different system. Although the suggestion of a separate controller thus exists in ForCES, the emphasis is on the communication between CE and FE over a switch backplane, as shown in Fig. 3.4.

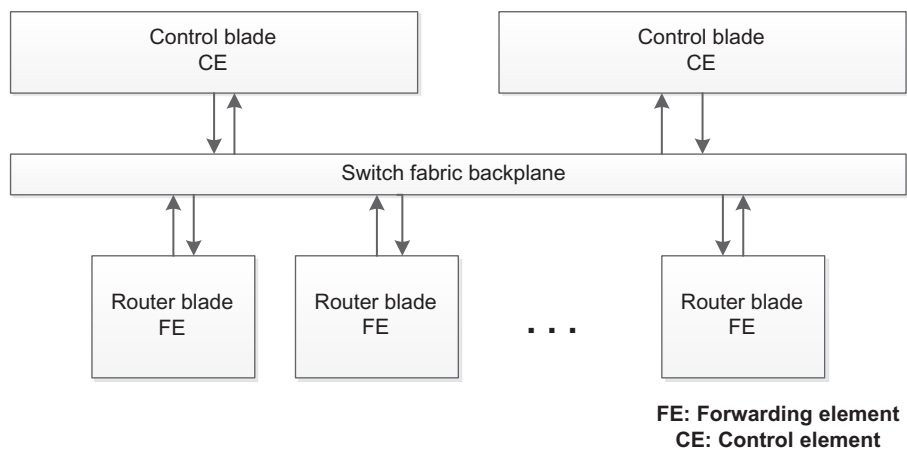


FIG. 3.4

ForCES design.

3.2.6 4D: CENTRALIZED NETWORK CONTROL

Seminal work on the topic of moving networking technology from distributed networking elements into a centralized controller appeared in the 4D proposal [20], *A Clean Slate 4D Approach to Network Control and Management*. 4D, named after the architecture's four planes—*decision*, *dissemination*, *discovery*, and *data*—proposes a complete refactoring of networking away from autonomous devices and toward the idea of concentrating control plane operation in a separate and independent system dedicated to that purpose.

4D argues that the state of networking today is fragile and, therefore, often teeters on the edge of failure because of its current design based on distributed, autonomous systems. Such systems exhibit a defining characteristic of unstable, complex systems: a small local event such as a misconfiguration of a routing protocol can have a severe, global impact. The proposal argues that the root cause is the fact that the control plane is running on the network elements themselves.

4D centers around three design principles:

- **Network-level Objectives:** In short, the goals and objectives for the network system should be stated in network-level terms based on the entire network domain, separate from the network elements, rather than in terms related to individual network device performance.
- **Network-wide View:** There should be a comprehensive understanding of the whole network. Topology, traffic, and events from the entire system should form the basis on which decisions are made and actions are taken.
- **Direct Control:** The control and management systems should be able to exert direct control over the networking elements, with the ability to program the forwarding tables for each device, rather than only being able to manipulate some remote and individual configuration parameters, as is the case today.

Fig. 3.5 shows the general architecture of a 4D solution, with centralized network control via the control and management system.

One aspect of 4D that is actually stated in the title of the paper is the concept of a *clean slate*, meaning the abandoning of the current manner of networking in favor of this new method as described by the three aforementioned principles. A quote from the 4D proposal states that “*We hope that exploring an extreme design point (the clean slate approach) will help focus the attention of the research and industrial communities on this crucially important and intellectually challenging area.*”

The 4D proposal delineates some of the challenges faced by the centralized architecture proposed. These challenges continue to be relevant today in SDN. We list a few of them here:

- **Latency:** Having a centralized controller means that a certain (hopefully small) number of decisions will suffer nontrivial round-trip latency as the networking element requests policy directions from the controller. How this delay impacts the operation of the network, and to what extent, remains to be determined. Also, with the central controller providing policy advice for a number of network devices, it is unknown whether the conventional servers on which the controller runs will be able to service these requests at sufficient speed to have minimal or no impact on network operation.
- **Scale:** Having a centralized controller means that responsibility for the topological organization of the network, determination of optimal paths, and responses to changes, must be handled by the controller. As has been argued, this is the appropriate location for this functionality; however, as more and more network devices are added to the network, questions arise of scale and the ability of a single controller to handle all those devices. It is difficult to know how well a centralized system

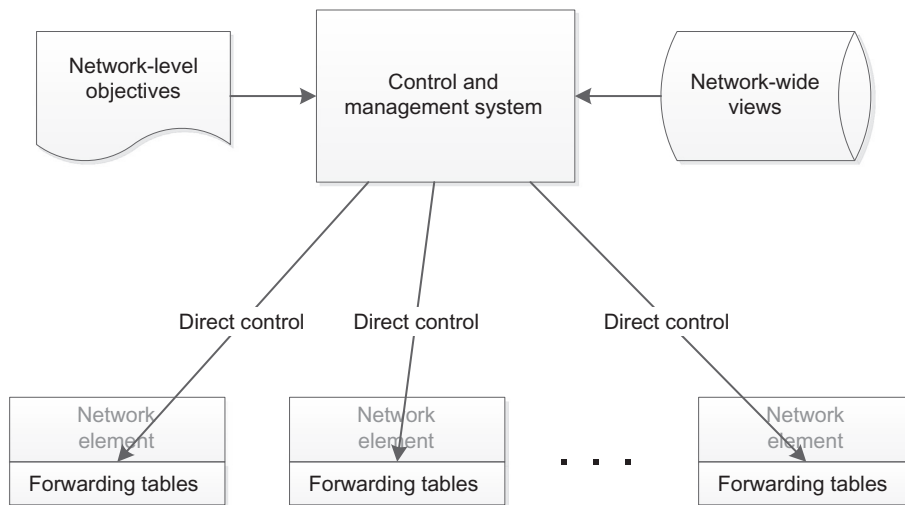


FIG. 3.5

4D principles.

can handle hundreds, thousands, or tens of thousands of network devices, nor what is the solution when the number of network devices outgrows the capacity of the controller to handle them. If we attempt to scale by adding additional controllers, how do they communicate, and who orchestrates coordination amongst the controllers? [Section 6.1.3](#) will address these questions.

- **High Availability (HA):** The centralized controller must not constitute a *single point of failure* for the network. This implies the need for redundancy schemes in a number of areas. First, there must be redundant controllers such that processing power is available in the event of failure of a single controller. Secondly, the actual data used by the set of controllers needs to be mirrored such that they can program the network devices in a consistent fashion. Thirdly, the communication paths to the different controllers need to be redundant to ensure that there is always a functioning communications path between a switch and at least one controller. We further discuss high availability in the context of a modern SDN network in [Section 6.1.2](#).
- **Security:** Having a centralized controller means that security attacks are able to focus on that one point of failure, and, thus, the possibility exists that this type of solution is more vulnerable to attack than a more distributed system. It is important to consider what extra steps must be taken to protect both the centralized controller and the communication channels between it and the networking devices.

ForCES and 4D contribute greatly toward the evolution of the concepts that underlie SDN—separation of forwarding and control planes (ForCES), and having a centralized controller responsible for overall routing and forwarding decisions (4D). However, both of these proposals suffer from a lack of actual implementations. Ethane, examined below, benefited from the learnings that can only come from a real-life implementation.

3.2.7 ETHANE: CONTROLLER-BASED NETWORK POLICY

Ethane was introduced in a paper entitled *Rethinking Enterprise Network Control* [21] in 2007. Ethane is a policy-based solution which allows network administrators to define policies pertaining to network-level access for users, which includes authentication and quarantine for misbehaving users. Ethane was taken beyond the proposal phase. Multiple instances have been implemented and shown to behave as suggested in [21].

Ethane is built around three fundamental principles:

- **The network should be governed by high-level policies:** Similar to 4D, Ethane espouses the idea that the network be governed by policies defined at high levels, rather than on a per-device basis. These policies should be at the level of services and users and the machines through which users can connect to the network.
- **Routing for the network should be aware of these policies:** Paths that packets take through the network are to be dictated by the higher-level policies described in the previous bullet point, rather than as in the case of today's networks, in which paths are chosen based on lower-level directives. For example, *guest* packets may be required to pass through a filter of some sort, or certain types of traffic may require routing across lightly loaded paths. Some traffic may be highly sensitive to packet loss, while other traffic (e.g., *Voice over IP* (VoIP)) may tolerate dropped packets but not latency and delay. These higher-level policies are more powerful guidelines for organizing and directing traffic flows than low-level and device-specific rules.
- **The network should enforce binding between packets and their origin:** If policy decisions rely on higher-level concepts, such as the concept of a *user*, then the packets circulating in the network must be traceable back to their point of origin (i.e., the user or machine which is the actual source of the packet).

Fig. 3.6 illustrates the basics of the Ethane solution. As will become apparent in the following chapters, there are many similarities between this solution and OpenFlow.

In order to test Ethane, the researchers themselves had to develop switches in order to have them implement the protocol and the behavior of such a device. That is, a device that allows control plane functionality to be determined by an external entity (the controller), and which communicates with that external entity via a protocol that allows flow entries to be configured into its local flow tables, which then perform the forwarding functions as packets arrive.

The behavior of the Ethane switches is generally the same as OpenFlow switches today which forward and filter packets based on the flow tables that have been configured on the device. If the switch does not know what to do with the packet, it forwards it to the controller and awaits further instructions.

In short, Ethane is basically a software defined networking technology, and its components are the antecedents of OpenFlow, which we describe in detail in [Chapter 5](#).

DISCUSSION QUESTION:

A key aspect of SDN is segregation of the control plane to a centralized device. Which of the technologies listed in [Table 3.1](#) used this approach?

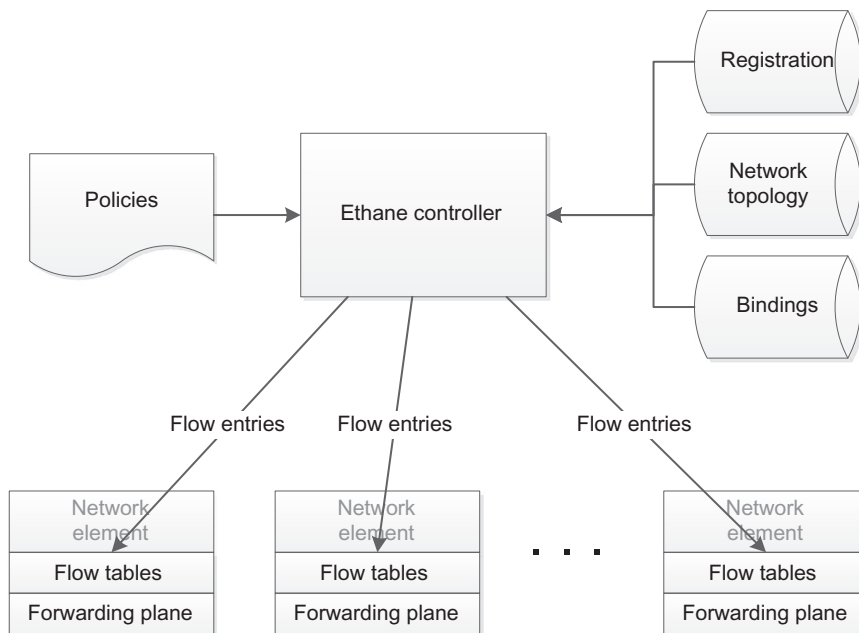


FIG. 3.6

Ethane architecture.

3.3 LEGACY MECHANISMS EVOLVE TOWARD SDN

The urgent need for SDN we described in [Chapter 2](#) could not wait for a complete new network paradigm to be fleshed out through years of research and experimentation. It is not surprising, then, that there were early attempts to achieve some SDN-like functionality within the traditional networking model. One such example was Cisco's *Policy Based Routing* that we described in [Chapter 1](#). The capabilities of legacy switches were sometimes extended to support detailed policy configuration related to security, QoS and other areas. Old APIs were extended to allow centralized programming of these features. Some SDN providers have based their entire SDN solution on a rich family of extended APIs on legacy switches, orchestrated by a centralized controller. In [Chapter 6](#) we will examine how these alternative solutions work and, whether or not they genuinely constitute an SDN solution.

3.4 SOFTWARE DEFINED NETWORKING IS BORN

3.4.1 THE BIRTH OF OPENFLOW

OpenFlow is a protocol specification that describes the communication between OpenFlow switches and an OpenFlow controller. Just as the previous sections presented standards and proposals which were precursors to SDN, seeing SDN through a gestation period, then the arrival of OpenFlow is the point at which SDN was actually born. In reality, the term SDN did not come into use until a year after OpenFlow made its appearance on the scene in 2008, but the existence and adoption of OpenFlow by

research communities and networking vendors marked a sea change in networking, one that we are still witnessing even now. Indeed, while the term SDN was in use in the research community as early as 2009, SDN did not begin to make a big impact in the broader networking industry until 2011.

For reasons identified in the previous chapter, OpenFlow was developed and designed to allow researchers to experiment and innovate with new protocols in everyday networks. The OpenFlow specification encouraged vendors to implement and enable OpenFlow in their switching products for deployment in college campus networks. Many network vendors have implemented OpenFlow in their products.

The OpenFlow specification delineates both the protocol to be used between the controller and the switch as well as the behavior expected of the switch. Fig. 3.7 illustrates the simple architecture of an OpenFlow solution.

The following list describes the basic operation of an OpenFlow solution:

- The controller populates the switch with flow table entries.
- The switch evaluates the header(s) of incoming packets and finds a matching flow, then performs the associated action. Depending on the match criteria, this evaluation would begin with the layer two header, and then potentially continue to the layer three and even layer four headers in some cases.
- If no match is found, the switch forwards the packet to the controller for instructions on how to deal with the packet.
- Typically the controller will update the switch with new flow entries as new packet patterns are received, so that the switch can deal with them locally. It is also possible that the controller will program *wildcard rules* that will govern many flows at once.




FIG. 3.7

General OpenFlow design.

OpenFlow will be examined in detail in [Chapter 5](#). For now, though, the reader should understand that OpenFlow has been adopted by both the research community and by a number of networking vendors. This has resulted in a significant number of network devices supporting OpenFlow on which researchers can experiment and test new ideas.

DISCUSSION QUESTION:

What similarities exist between Ethane and OpenFlow?

3.4.2 OPEN NETWORKING FOUNDATION

OpenFlow began with the publication of the original proposal in 2008. By 2011 OpenFlow had gathered enough momentum that the responsibility for the standard itself moved to the *Open Networking Foundation* (ONF). The ONF was established in 2011 by Deutsche Telekom, Facebook, Google, Microsoft, Verizon, and Yahoo!. It is now the guardian of the OpenFlow standard, and consists of a number of councils, areas and working groups. We depict the interrelationships between these bodies in [Fig. 3.8 \[27\]](#).

One novel aspect of the ONF is that corporate members of the Board of Directors consist of major network *operators*, and not the networking *vendors* themselves. As of this writing, the ONF board is composed of CTOs, Technical Directors, professors from Stanford and Princeton, and Fellows from companies such as Google, Yahoo!, Facebook, Deutsche Telekom, Verizon, Microsoft, NTT, and Goldman Sachs among others. This helps to prevent the ONF from supporting the interests of one major networking vendor over another. It also helps to provide a *real-world* perspective on what should be investigated and standardized. Conversely, it runs the risk of defining a specification that is difficult for NEMs to implement. Our previous comments about NEMs being locked into their status quo notwithstanding, it is also true that there is a wealth of engineering experience resident in the NEMs regarding how to actually design and build high performance, high reliability switches. While the real-world experience of the users is indeed indispensable, it is imperative that the ONF seek input from the vendors to ensure that the specifications they produce are in fact implementable.

3.5 SUSTAINING SDN INTEROPERABILITY

At the current point in the evolution of SDN, we have a standard which has been accepted and adopted by academia and industry alike, and we have an independent standards body to shepherd OpenFlow forward and to ensure it is independent and untethered to any specific institution or organization. It is now important to ensure that the implementations of the various players in the OpenFlow space adhere to the standards as they are defined, clarify and expand the standards where they are found to be incomplete or imprecise, and in general guarantee interoperability amongst OpenFlow implementations. This goal can be achieved in a few ways:

- **Plugfests:** Plugfests, staged normally at conferences, summits, and congresses, are environments where vendors can bring their devices and software in order to test them with devices and software

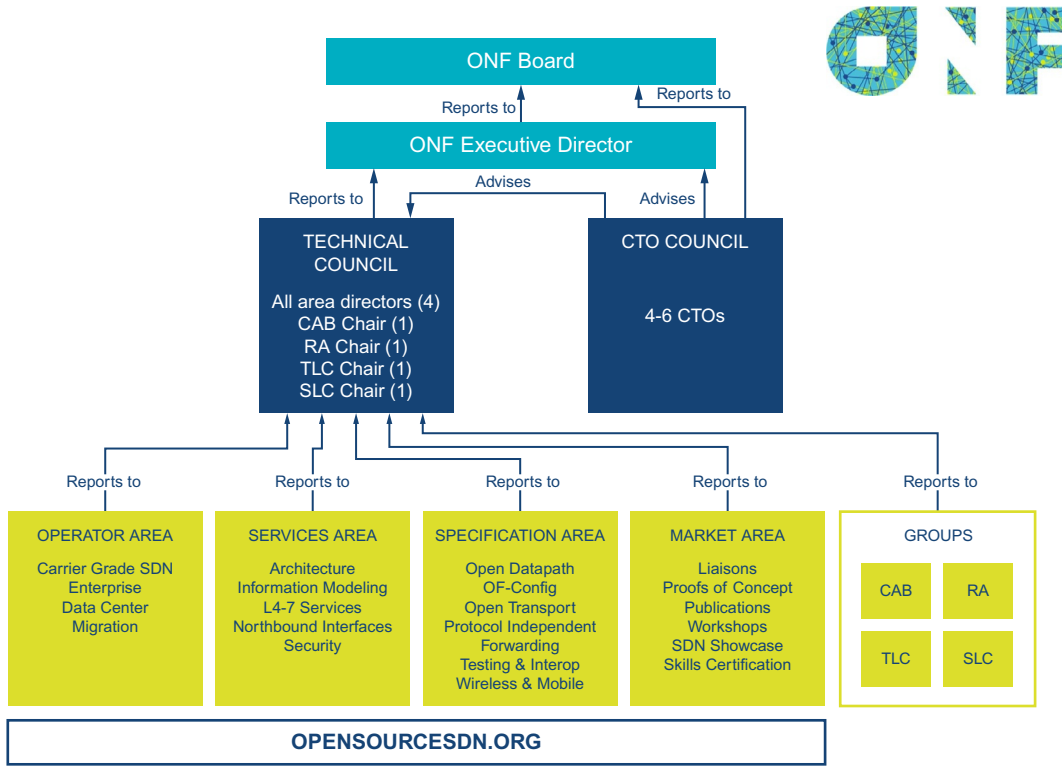


FIG. 3.8

ONF orgchart.

Reproduced with permission from

<https://www.opennetworking.org/certification/158-module-content/technical-communities-modules/1908-org-chart>.

from other vendors. These are rich opportunities to determine where implementations may be lacking, or where the standard itself is unclear and needs to be made more precise and specific.

- **Interoperability Labs:** Certain institutions have built dedicated test labs for the purpose of testing the interoperability of equipment from various vendors and organizations. One such lab is the *Indiana Center for Network Translational Research and Education* (InCNTRE) at Indiana University, which hosts a large collection of vendor devices and controllers, as well as experimental devices and controllers from open source contributors. We discuss open source contributions to SDN in the next section.
- **Certification Programs:** There is a need for certification of switches so buyers can know they are getting a switch that is certified to support a particular version(s) of OpenFlow. The ONF has now implemented such a program [22].
- **Education and Consulting:** A complex, game-changing technological shift such as that represented by SDN will not easily permeate a large industry without the existence of an infrastructure to train and advise networking staff about the migration. It is important that a cadre of highly qualified yet vendor-neutral organizations address this need.

Initially, many SDN interoperability tests revealed dissimilarities and issues. While existing implementations of OpenFlow are increasingly interoperable, challenges remain. For example, as of this writing, it remains difficult for an OpenFlow controller to work consistently across multiple, varied switch implementations of OpenFlow. This is largely due to limitations in existing ASICs that lead to switches only supporting different subsets of OpenFlow consistent with their particular hardware. Other problems emanate from the fact that the OpenFlow 1.0 specification is vague in terms of specifying which features are required versus optional. Later versions of the specification attempt to clarify this, but even the latest specifications suffer from a large number of optional features for which support is far from universal. As we will see in [Chapter 5](#), the OpenFlow standard is not static, and the goalposts of interoperability are moved with each new release.

3.6 OPEN SOURCE CONTRIBUTIONS

One of the basic rationales for SDN is that innovation and advancement have been stifled as a result of the closed and inflexible environment that exists in networking today. Thus, the openness that results from the creation of standards such as OpenFlow should encourage researchers to dissect old networking methods, and should usher in a new dawn of network operation, management, and control. This section will examine ways in which open source contributes to this process.

3.6.1 THE POWER OF THE COLLECTIVE

Technological advancement sometimes arises from the efforts of corporations and major organizations, quite often due to the fact that they are the only ones in the position to make contributions in their domains. In the world of software, however, it is occasionally possible for small players to develop technology and make it freely available to the general public. Some examples are:

- **Operating Systems:** The Linux operating system was developed as open source and is used today to control countless devices that we use every day, from *Digital Video Recorders* (DVRs) to smartphones.
- **Databases:** Many of the websites we visit for news reports or to purchase products store their information and product data in databases that have been developed, at least initially, by the open source community. MySQL is an example of such an open source data base.
- **Servers:** When we access locations on the Internet, many of those servers that we are accessing are running application server software and using tools which have been developed over time by the open source community. The open source Apache Web Server is used in countless applications worldwide.
- **Security:** Applying the open source model is also often considered to deliver more secure environments [23]. Open source can be more secure because of the peer review and white box evaluation that naturally occur in the open source development paradigm. Proprietary protocols may be less secure because they are not open and evaluated. Many security products providing antivirus protection and maintaining lists of malicious sites and programs are running open source software to accomplish their tasks. OpenSSL is probably the foremost example of a widely used open source encryption toolkit.

- **File Sharing:** The BitTorrent protocol is an example of a hugely successful [24] open protocol used for file sharing. BitTorrent works as a *P2P/Overlay* network that achieves high-bandwidth file downloading by performing the download of a file piecemeal and in parallel from multiple servers.

These are just a few examples. Imagine our world today without those items listed above. Would private institutions have eventually implemented the software solutions required in order to provide that functionality? Most likely. Would these advancements in technology have occurred at the velocity that we have witnessed in the past ten years, without current and ongoing contributions from open source? Almost certainly not.

3.6.2 THE DANGER OF THE COLLECTIVE

Of course, with an endeavor being driven by individuals who are governed not only by their own desire to contribute but also by their free will, whims and other interests, there is bound to be some risk. The areas of risk include quality, security, timeliness, and support. We explain below how these risks can be mitigated.

Open source software must undergo tests and scrutiny by even larger numbers of individuals than its commercial counterpart. This is due to the fact that an entire world of individuals has access to and can test those contributions. For any given feature being added or problem being solved, the open source community may offer a number of competing approaches to solve the problem. Even Open Source initiatives are subject to release cycles and there are key individuals involved in deciding what code will make its way into the next release. Just because an open source developer creates a body of code for an open source product does not mean that it makes it into a release. The fact that competing approaches may be assessed by the community and that admission into a release is actually controlled by key individuals associated with the open source effort manage to keep quality high. These same factors serve to minimize the risk of security threats. Since the source code is open for all to view, it is more difficult to hide malicious threats such as back doors.

The issue of timeliness is more nebulous. Unlike a project sponsored and fully funded by one's own organization, only indirect influence can be applied to exactly when a particular feature appears in an open source product. The prudent approach is not to depend on future deliverables, but to use existing functionality.

There seem to be two fundamental solutions to the issue of support. If you are a business intending to utilize open source in a product you are developing, you either need to feel that you have the resources to support it on your own, or that there is such a large community that has been using that code for a long enough time that you simply trust that the bad bugs have already surfaced and that you are using a stable code base.

It is important to remember that open source licensing models differ greatly from one model to the next. Some models severely restrict how contributions can be made and are used. We discuss in more detail in [Section 13.4](#) some of the different major open source licenses and the issues with them.

3.6.3 OPEN SOURCE CONTRIBUTIONS TO SDN

Based on the previous discussions, it is easy to see the potential value of open source contributions in the drive toward SDN. Huge advances in SDN technology are attributable to open source projects.

Multiple open source implementations of SDN switches, controllers and applications are available. In [Chapter 13](#) we provide details of the open source projects that have been specifically targeted to accelerate innovation in SDN. In that chapter we will also discuss other open source efforts which, while not as directly related to SDN, are nonetheless influential on the ever-growing acceptance of the SDN paradigm.

3.7 NETWORK VIRTUALIZATION

In [Chapter 2](#) we discussed how *network virtualization* lagged behind its compute and storage counterparts and how this has resulted in a strong demand for network virtualization in the data center. Network virtualization, in essence, provides a network service that is decoupled from the physical hardware below that offers a feature set identical to the behavior of its physical counterpart. An important and early approach to such network virtualization was the *Virtual Local Area Network* (VLAN). VLANs permitted multiple virtual local area networks to co-reside on the same layer two physical network in total isolation from one another. While this technical concept is very sound, the provisioning of VLANs is not particularly dynamic, and they only scale to the extent of a layer two topology. Layer three counterparts based on tunneling scale better than VLANs to larger topologies. Complex systems have evolved to use both VLAN as well as tunneling technologies to provide network virtualization solutions.

One of the most successful commercial endeavors in this space was Nicira, now part of VMware. Early on, Nicira claimed that there were seven properties of network virtualization [\[25\]](#):

1. Independence from network hardware
2. Faithful reproduction of the physical network service model
3. Follow operational model of compute virtualization
4. Compatibility with any hypervisor platform
5. Secure isolation between virtual networks, the physical networks, and the control plane
6. Cloud performance and scale
7. Programmatic network provisioning and control

Several of these characteristics closely resemble what we have said is required from an SDN solution. SDN promises to provide a mechanism for automating the network and abstracting the physical hardware below from the software defined network above. Network virtualization for data centers has undoubtedly been the largest commercial driver behind SDN. This momentum has become so strong that to some network virtualization has become synonymous with SDN. Indeed, VMware's (Nicira) standpoint [\[26\]](#) on this issue is that SDN is simply about abstracting control plane from data plane, and, therefore, network virtualization is SDN. Well, is it SDN or not?

3.8 MAY I PLEASE CALL MY NETWORK SDN?

If one were to ask four different attendees at a networking conference in 2016 what they thought qualified a network to be called SDN, they would likely provide divergent answers. Based on the genesis of SDN as presented in [Section 3.4](#), in this book we will define an SDN network as characterized by five

fundamental traits: *plane separation*, *a simplified device*, *centralized control*, *network automation and virtualization*, and *openness*. We will call an SDN solution possessing these five traits an *Open SDN* technology. We acknowledge that there are many competing technologies offered today that claim that their solution is an SDN solution. Some of these technologies have had larger economic impact in terms of the real-life deployments and dollars spent by customers than those that meet all five of our criteria. In some respects they may address customers' needs better than Open SDN. For example, a network virtualization vendor such as Nicira has had huge economic success and widespread installations in data centers, but does this without simplified devices. We define these five essential criteria not to pass judgment on these other SDN solutions, but in acknowledgement of what the SDN pioneers had in mind when they coined the term SDN in 2009 to refer to their work on OpenFlow. We will provide details about each of these five fundamental traits in [Section 4.1](#), and compare and contrast competing SDN technologies against these five as well as other criteria in [Chapter 6](#).

DISCUSSION QUESTION:

Looking back at [Table 3.1](#), after Ethane, which two of the technologies listed there most closely approach what we have defined in this book as SDN? Why?

3.9 CONCLUSION

With the research and open source communities clamoring for an open environment for expanded research and experimentation, and the urgent needs of data centers for increased agility and virtualization, networking vendors have been forced into the SDN world. Some have moved readily into that world, while others have dragged their feet or have attempted to redefine SDN. In the next chapter we will examine what SDN is and how it actually works.

REFERENCES

- [1] Ferro G. Networking vendors should step up with an SDN strategy. Network Computing, June 7, 2012. Retrieved from: <http://www.networkcomputing.com/next-gen-network-tech-center/networking-vendors-should-step-up-with-a/240001600>.
- [2] Godbole A. Data communications and networks. New Delhi: Tata McGraw-Hill; 2002.
- [3] Mendonca M, Astuto B, Nunes A, Nguyen X, Obraczka K, Turletti T. A survey of software-defined networking: past, present and future of programmable networks. IEEE Commun Surv Tutor 2014;16(3):1617–34.
- [4] Feamster N, Rexford J, Zegura E. The road to SDN: an intellectual history of programmable networks. ACM Mag 2013;11(12).
- [5] Devolved control of ATM networks. Retrieved from: <http://www.cl.cam.ac.uk/research/srg/netos/old-projects/dcan/>.
- [6] Campbell A, Katzela I, Miki K, Vicente J. Open signalling for ATM, Internet and Mobile Networks (OPENSIG'98). ACM SIGCOMM Comput Commun Rev 1999;29(1):97–108.
- [7] Doria A, Hellstrand F, Sundell K, Worster T. General Switch Management Protocol (GSMP) V3, RFC 3292. Internet Engineering Task Force; 2002.

- [8] A comparison of IP switching technologies from 3Com, Cascade, and IBM. CSE 588 Network Systems, Spring, 1997. University of Washington; 1997. Retrieved from: <http://www.cs.washington.edu/education/courses/csep561/97sp/paper1/paper11.txt>.
- [9] Tennehouse D, Smith J, Sincoskie W, Wetherall D, Minden, G. A survey of active network research. *IEEE Commun Mag* 1997;35(1):80–6.
- [10] Tennehouse D, Wetherall D. Towards an active network architecture. In: *Proceedings of the DARPA Active networks conference and exposition, 2002, IEEE; 2002*. p. 2–15.
- [11] Bhattacharjee S, Calvert K, Zegura E. An architecture for active networking. In: *High performance networking VII: IFIP TC6 Seventh international conference on High Performance Networks (HPN '97), April–May 1997, White Plains, NY, USA; 1997*.
- [12] Schwartz B, Jackson A, Strayer W, Zhou W, Rockwell R, Partridge C. Smart packets for active networks. In: *Open architectures and network programming proceedings, IEEE OPENARCH '99, March 1999, New York, NY, USA; 1999*.
- [13] da Silva S, Yemini Y, Florissi D. The NetScript active network system. *IEEE J Sel Areas Commun* 2001;19(3):538–51.
- [14] Rigney C, Willens S, Rubens A, Simpson W. Remote Authentication Dial In User Service (RADIUS), RFC 2865. Internet Engineering Task Force; 2000.
- [15] Durham D, Boyle J, Cohen R, Herzog S, Rajan R, Sastry A. The COPS (Common Open Policy Service) Protocol, RFC 2748. Internet Engineering Task Force; 2000.
- [16] Ferro G. Automation and orchestration. *Network Computing*, September 8, 2011. Retrieved from: <http://www.networkcomputing.com/private-cloud-tech-center/automation-and-orchestration/231600896>.
- [17] VMWare vCenter. Retrieved from: <http://www.vmware.com/products/vcenter-server/overview.html>.
- [18] VMWare vMotion. Retrieved from: <http://www.vmware.com/files/pdf/VMware-VMotion-DS-EN.pdf>.
- [19] Doria A, Hadi Salim J, Haas R, Khosravi H, Wang W, Dong L, et al. Forwarding and Control Element Separation (ForCES) protocol specification, RFC 5810. Internet Engineering Task Force; 2010.
- [20] Greenberg A, Hjalmtysson G, Maltz D, Myers A, Rexford J, Xie G, et al. A clean slate 4D approach to network control and management. *ACM SIGCOMM Comput Commun Rev* 2005;35(3).
- [21] Casado M, Freedman M, Pettit J, McKeown N, Shenker S. Ethane: taking control of the enterprise. *ACM SIGCOMM Comput Commun Rev* 2007;37(4):1–12.
- [22] Lightwave Staff. Open Networking Foundation launches OpenFlow certification program. Lightwave, July 2013. Retrieved from: <http://www.lightwaveonline.com/articles/2013/07/open-networking-foundation-launches-openflow-certification-program.html>.
- [23] Benefits of Open Source Software. Open Source for America. Retrieved from: <http://opensourceforamerica.org/learn-more/benefits-of-open-source-software/>.
- [24] BitTorrent and μ Torrent software surpass 150 million user milestone; announce new consumer electronics partnerships. BitTorrent, January 2012. Retrieved from: http://www.bittorrent.com/intl/es/company/about/ces_2012_150m_users.
- [25] Gourley B. The seven properties of network virtualization. CTOvision, August 2012. Retrieved from: <http://ctovision.com/2012/08/the-seven-properties-of-network-virtualization/>.
- [26] Metz C. What is a virtual network? It's not what you think it is. *Wired*, May 9, 2012. Retrieved from: <http://www.wired.com/wiredenterprise/2012/05/what-is-a-virtual-network>.
- [27] Org Chart. Open Networking Foundation; 2016. Retrieved from: <https://www.opennetworking.org/certification/158-module-content/technical-communities-modules/1908-org-chart>.